SRI CHANDRASEKHARENDRA SARASWATHI VISWA MAHAVIDYALAYA

(UNIVERSITY ESTABLISHED UNDER SECTION 3 OF UGC ACT 1956)
ENATHUR, KANCHIPURAM – 631 561

DEPARTMENT OF COMPUTER SCIENCE & ENGINEERING

Name : penta Bhargavi Sai Abhi Sri Ishwarya

Reg. No: 11249A28

Class : II B.E. (CSE)

Course Code: BCSF244A12

COURSE NAME: Data Science For Engineers

SRI CHANDRASEKHARENDRA SARASWATHI VISWA MAHAVIDYALAYA

(UNIVERSITY ESTABLISHED UNDER SECTION 3 OF UGC ACT 1956)
ENATHUR, KANCHIPURAM – 631 561

DEPARTMENT OF COMPUTER SCIENCE & ENGINEERING

BONAFIDE CERTIFICATE

This is to certify that this is the bonafide record of work done by Mr/Ms. **Penta bhargavi sai abhi sri ishwarya**,

with Reg. No **11249A282** of **II-B.E.(CSE)** during the academic year 2025 – 2026.

Station:

Date:

Staff-in-charge **Head of the Department**

Submitted for the Practical examination held on_____.

Examiner-1 **Examiner-2**

INDEX

SL.NO	ProgramName	Date	Page no	Signature
1	Introduction to python libraries		4-10	
a	Numpy	29/01/26	4-5	
b	Pandas	29/01/26	5-6	
c	Matplotlib	05/02/26	6-8	
d	Scikit-learn	05/02/26	9-10	
2	Perform data exploration and preprocessing python	13/2/26	11-12	
3	Implement Regularized Linear Regression	19/2/26	13-13	
4	Implement Naive Bayes Classifier	5/3/26	14-14	
5	Implement Regularized Logistic Regression	5/3/26	15-16	
6	Build models using ensembling techniques	12/3/26	17-17	
7	Build models using decision tree	12/3/26	18-18	
8	Support vector machines with different Kernels	2/4/26	19-19	
9	Implement K-Nearest Neighbours Classification	2/4/26	20-20	

10	K-Means clustering with PCA and elbow method	2/4/26	21-22	

Exp 1: Introduction to python libraries

1.1 NumPy

Aim: To implement array creation and operations using NumPy.

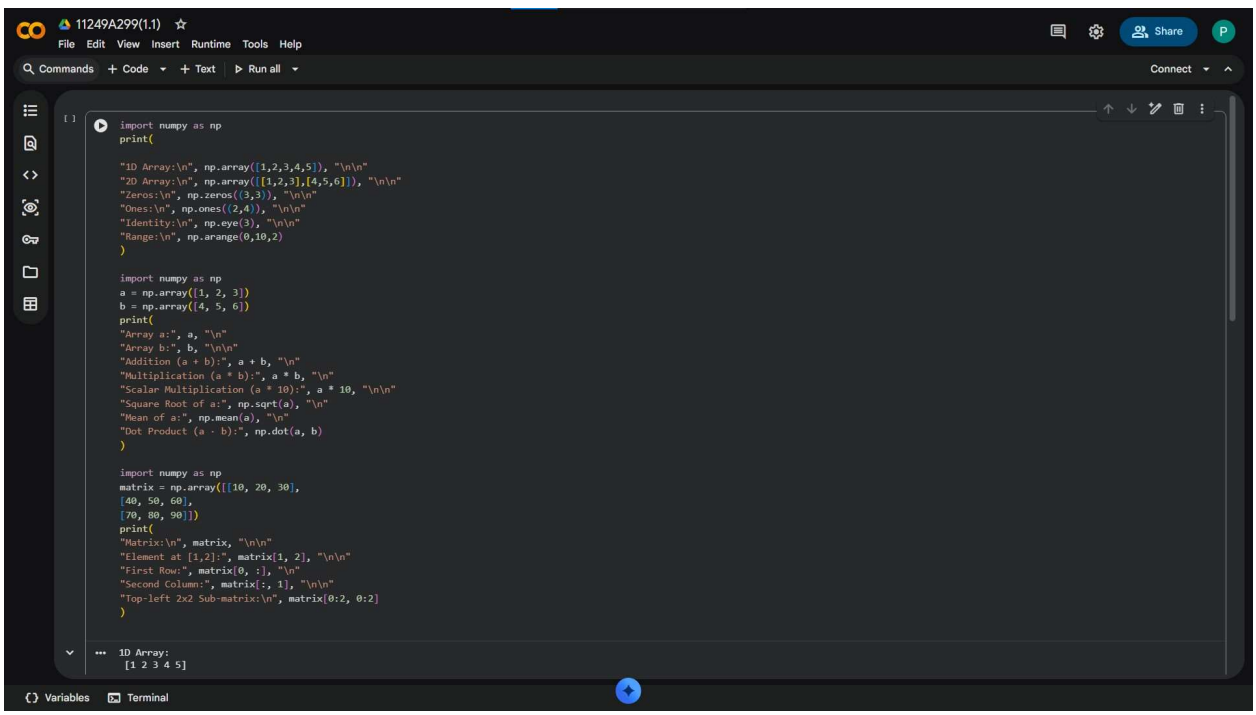

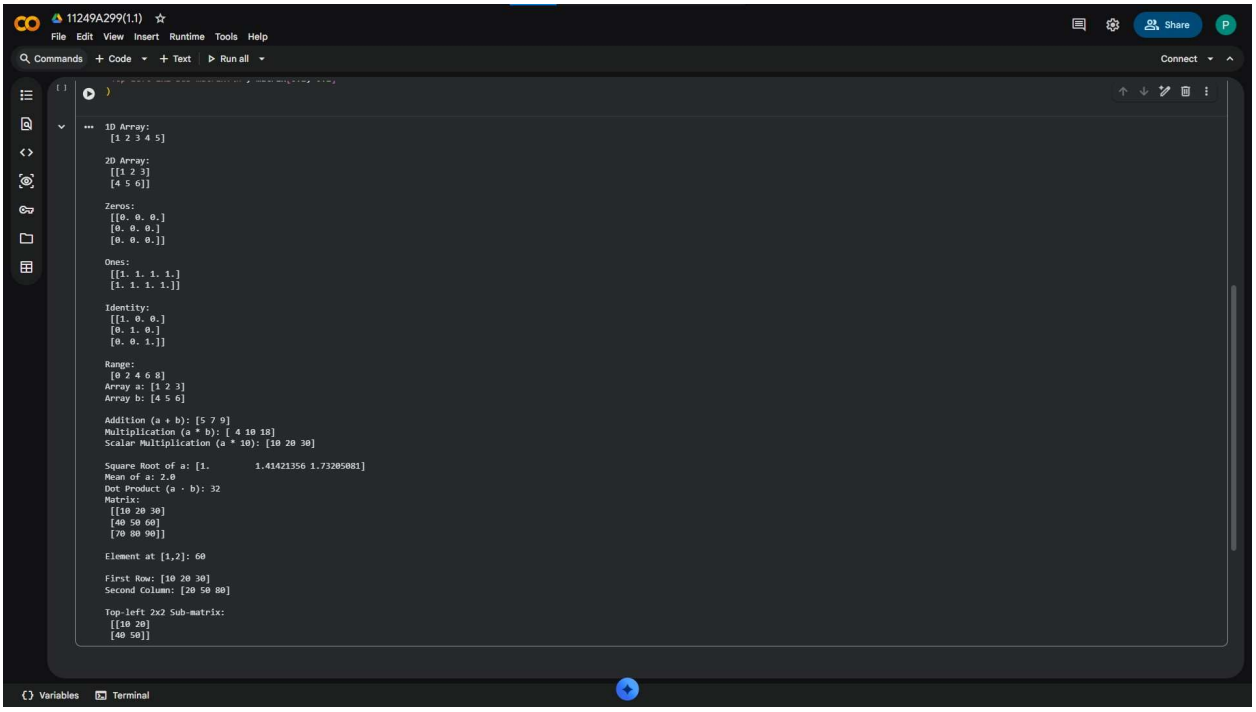

1.2 Pandas

Aim: To implement data creations, filtering, and cleaning using Pandas

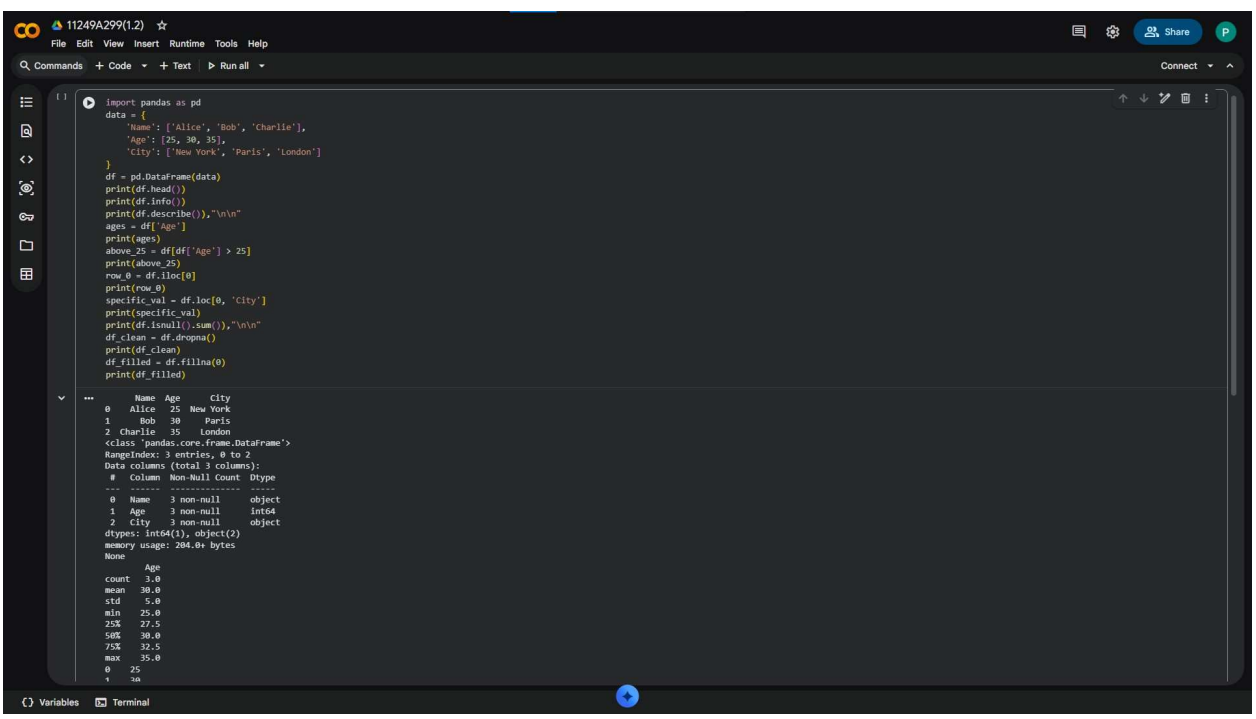

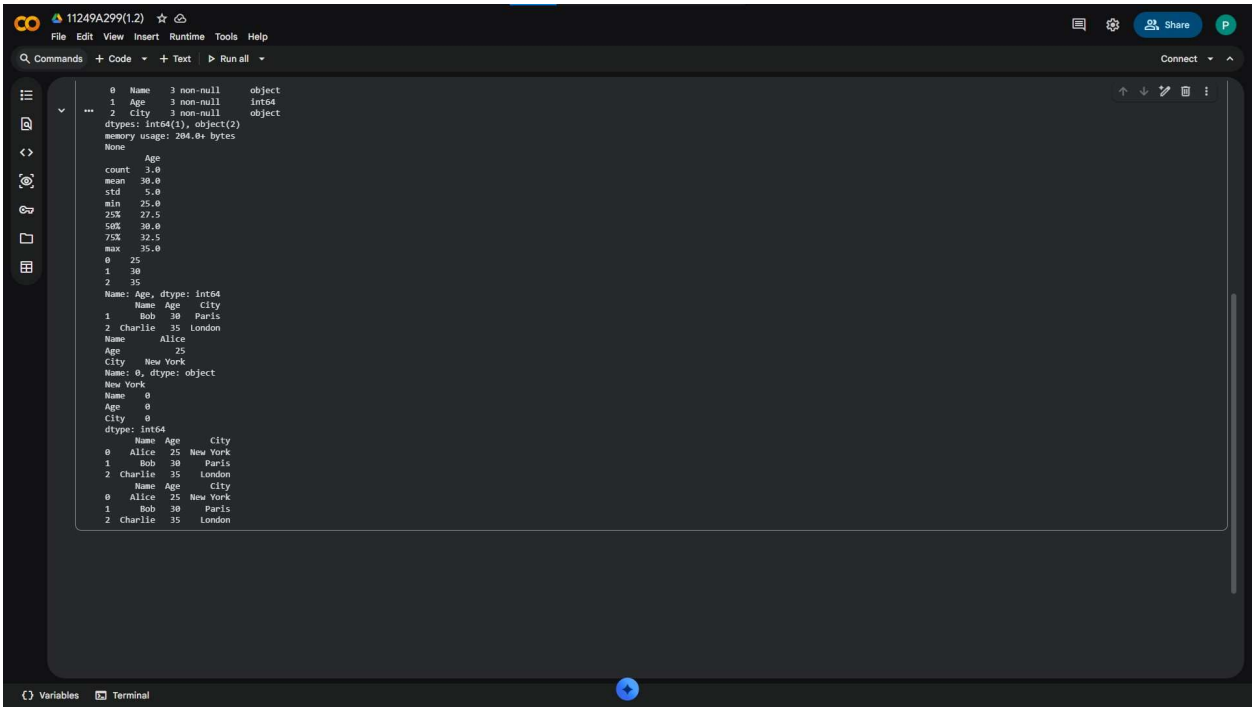

1.3 Matplot lib

Aim: To implement plotting operations using Matplolib.

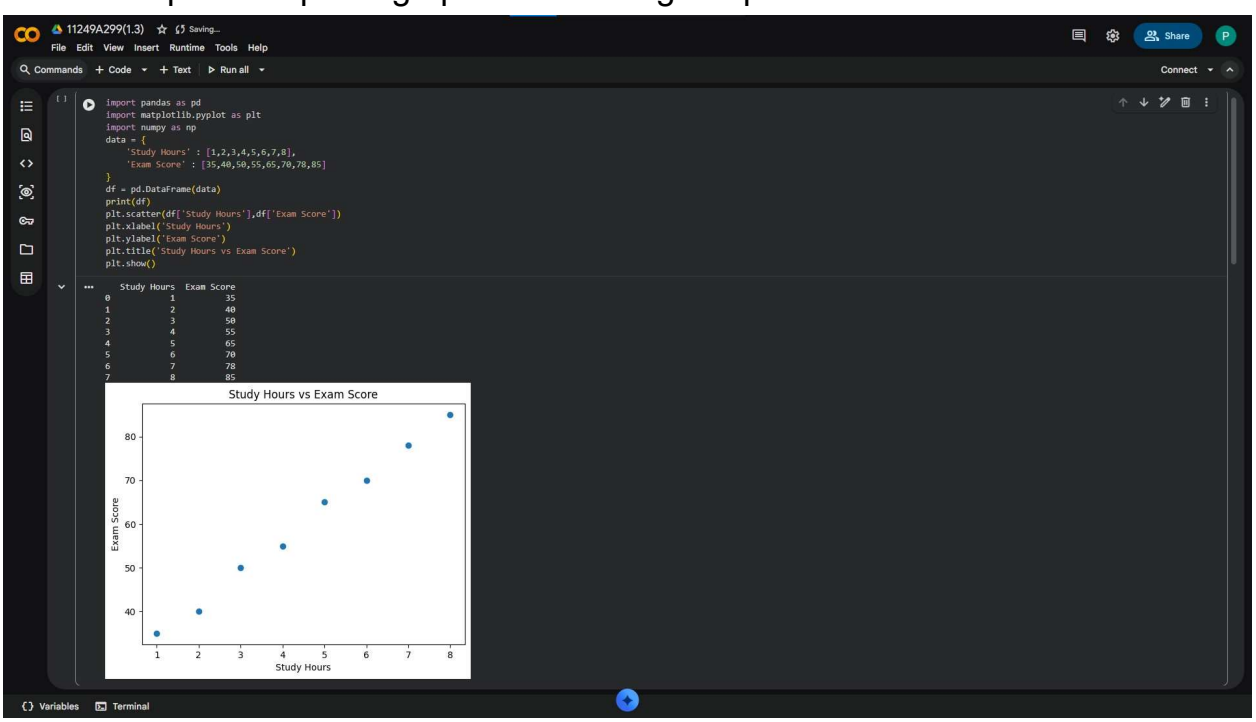

11249A299(1.3)

File Edit View Insert Runtime Tools Help

Commands + Code + Text Run all

Share Connect

```
plt.plot (df['Study Hours'],df['Exam Score'],marker='o')
plt.xlabel('Study Hours')
plt.ylabel('Exam Score')
plt.title('Exam Score Progress')
plt.show()
```

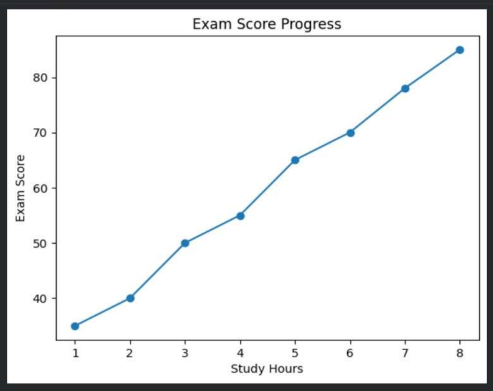

Variables Terminal

11249A299(1.3)

File Edit View Insert Runtime Tools Help

Commands + Code + Text Run all Connect Share

```
plt.hist(df['Exam Score'],bins=5)
plt.xlabel('Exam Score')
plt.ylabel('Frequency')
plt.title('Exam Score Distribution')
plt.show()
```

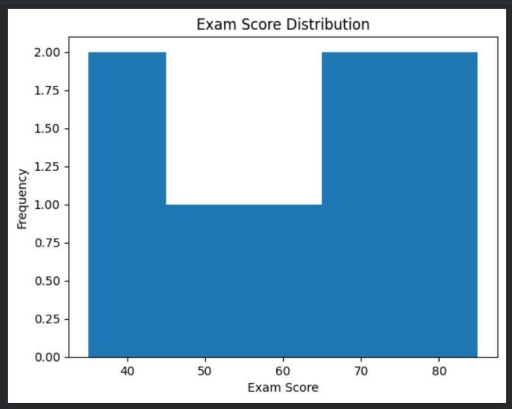

Variables Terminal

11249A299(1.3)

File Edit View Insert Runtime Tools Help

Commands + Code + Text Run all Connect Share

```
plt.bar (df['Study Hours'],df['Exam Score'])
plt.xlabel('Study Hours')
plt.ylabel('Exam Score')
plt.title('Study Hours vs Exam Score')
plt.show()
```

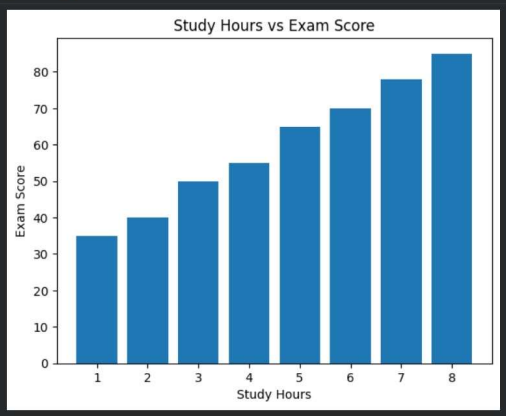

Variables Terminal

1.4 Scikit-Learn

Aim: To implement machine learning algorithms using Scikit-learn.

```python
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
data = {
    'Student' : ['S1','S2','S3','S4','S5','S6','S7','S8','S9','S10'],
    'classes_attended' : [30,35,40,45,50,55,60,65,70,75],
    'Internal_marks' : [35,38,42,45,47,50,52,55,58,60],

}
df = pd.DataFrame(data)
print(df)

X=df[['classes_attended']]
print(X)
y=df['Internal_marks']
print(y)

X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.2,random_state=7)
print(X_train)
print(X_test)
print(y_train)
print(y_test)

model = LinearRegression()
model.fit(X_train,y_train)

predictions = model.predict(X_test)

mse = mean_squared_error(y_test, predictions)
print("MSE:", mse)
print("Marks per Class:", model.coef_[0])
print("Base Marks:", model.intercept_)

plt.scatter(X, y)
plt.plot(X, model.predict(X))
plt.xlabel("Classes Attended")
plt.ylabel("Internal Marks")
```

```python
plt.scatter(X, y)
plt.plot(X, model.predict(X))
plt.xlabel("Classes Attended")
plt.ylabel("Internal Marks")
plt.title("Classes Attended vs Internal Marks")
plt.show()
```

```
  Student  classes_attended  Internal_marks
0      S1                30              35
1      S2                35              38
2      S3                40              42
3      S4                45              45
4      S5                50              47
5      S6                55              50
6      S7                60              52
7      S8                65              55
8      S9                70              58
9     S10                75              60
   classes_attended
0                30
1                35
2                40
3                45
4                50
5                55
6                60
7                65
8                70
9                75
0    35
1    38
2    42
3    45
4    47
5    50
6    52
7    55
8    58
9    60
Name: Internal_marks, dtype: int64
```

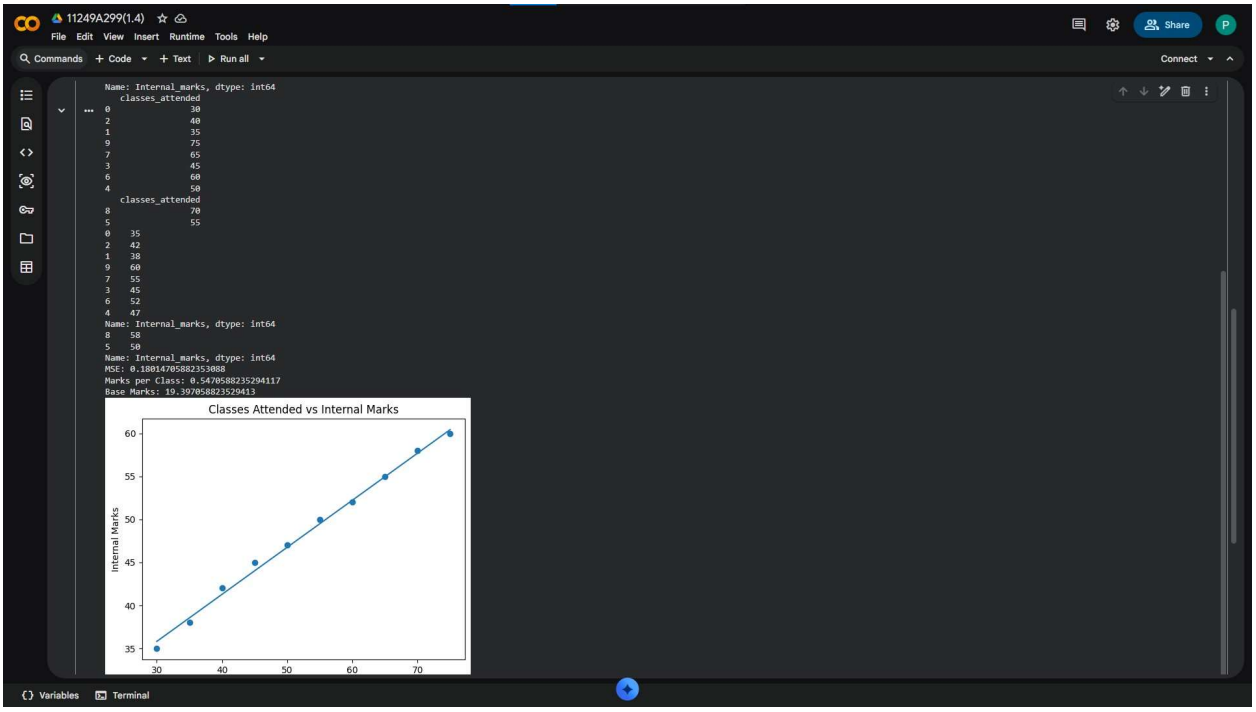

Result: Machine learning models were implemented successfully.

Exp 2 : Perform Data Exploration and preprocessing in python

Aim: To implement data exploration and preprocessing.

```python
import pandas as pd
import numpy as np
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split

data = {
'Country': ['India', 'USA', 'India', 'USA', 'UK', 'India'],
'Age': [22, 25, np.nan, 30, 28, 35],
'Salary': [40000, np.nan, 50000, np.nan, 72000, 58000],
'Purchased': ['No', 'Yes', 'Yes', 'No', 'Yes', 'Yes']
}
df = pd.DataFrame(data)
print("--- ORIGINAL RAW DATA ---")
print(df)
print("\n")

print("--- MISSING VALUES COUNT ---")
print(df.isnull().sum())
print("\n")
df['Age'] = df['Age'].fillna(df['Age'].mode())
df['Salary'] = df['Salary'].fillna(df['Salary'].mode())

print("--- DATA AFTER CLEANING ---")
print(df)
print("\n")
df_encoded = pd.get_dummies(df, columns=['Country'])
df_encoded['Purchased'] = df_encoded['Purchased'].map({'Yes': 1, 'No': 0})

print("--- DATA AFTER ENCODING (All Numbers Now) ---")
print(df_encoded)
print("\n")

X = df_encoded.drop('Purchased', axis=1)
y = df_encoded['Purchased']

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
print("--- FINAL PROCESSED DATA (Ready for AI) ---")
print(pd.DataFrame(X_scaled, columns=X.columns).head())
```

```
--- ORIGINAL RAW DATA ---
  Country   Age   Salary Purchased
0   India  22.0  40000.0       No
1     USA  25.0      NaN      Yes
2   India   NaN  50000.0      Yes
3     USA  30.0      NaN       No
4      UK  28.0  72000.0      Yes
```

```
3     USA  30.0      NaN       No
4      UK  28.0  72000.0      Yes
5   India  35.0  58000.0      Yes


--- MISSING VALUES COUNT ---
Country      0
Age          1
Salary       2
Purchased    0
dtype: int64


--- DATA AFTER CLEANING ---
  Country   Age   Salary Purchased
0   India  22.0  40000.0       No
1     USA  25.0  50000.0      Yes
2   India  28.0  50000.0      Yes
3     USA  30.0  72000.0       No
4      UK  28.0  72000.0      Yes
5   India  35.0  58000.0      Yes


--- DATA AFTER ENCODING (All Numbers Now) ---
   Age   Salary  Purchased  Country_India  Country_UK  Country_USA
0 22.0  40000.0          0           True       False        False
1 25.0  50000.0          1          False       False         True
2 28.0  50000.0          1           True       False        False
3 30.0  72000.0          0          False       False         True
4 28.0  72000.0          1          False        True        False
5 35.0  58000.0          1           True       False        False


--- FINAL PROCESSED DATA (Ready for AI) ---
        Age    Salary  Country_India  Country_UK  Country_USA
0 -1.484615 -1.438476           1.0   -0.447214    -0.707107
1 -0.742307 -0.592314          -1.0   -0.447214     1.414214
2  0.000000 -0.592314           1.0   -0.447214    -0.707107
3  0.494872  1.269243          -1.0   -0.447214     1.414214
4  0.000000  1.269243          -1.0    2.236068    -0.707107
```

Result : Data preprocessing was performed successfully.

Exp

3 : Implement Regularized Linear Regression

Aim: To implement regularized linear regression.

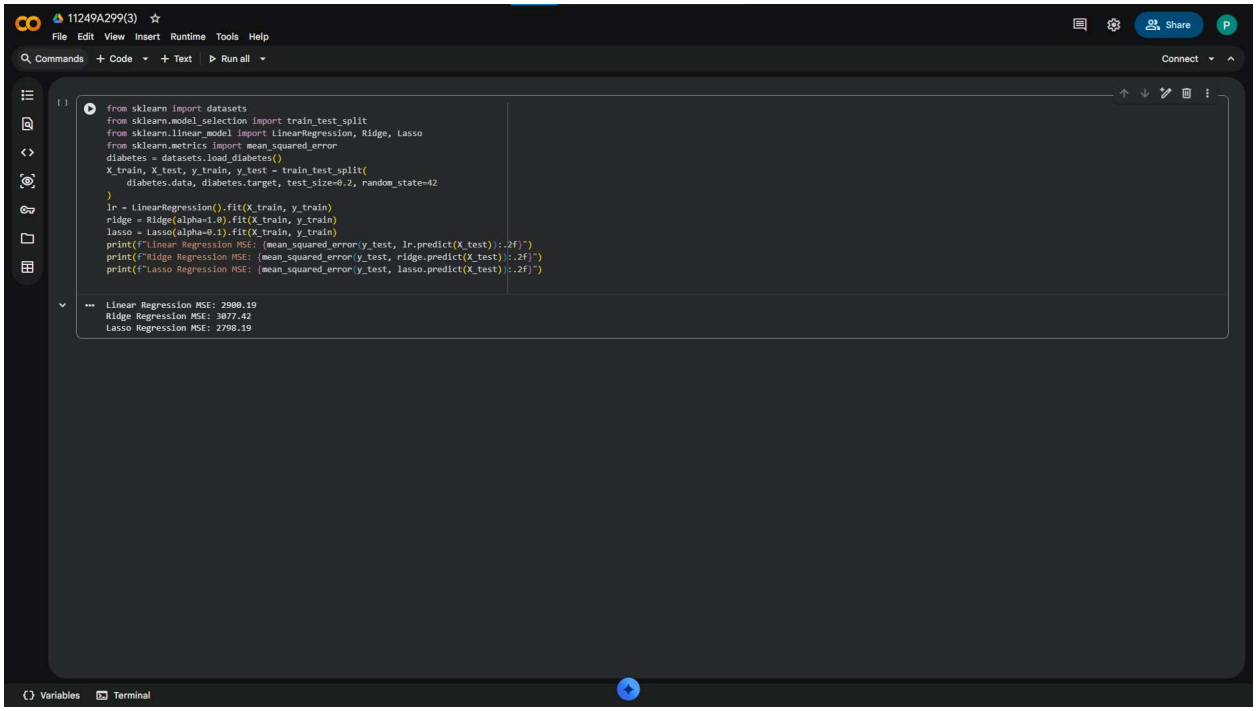

Result : The regression model was implemented successfully

Exp 4 : Implement Naive Bayes Classifier

Aim: To implement Naive Bayes Classification.

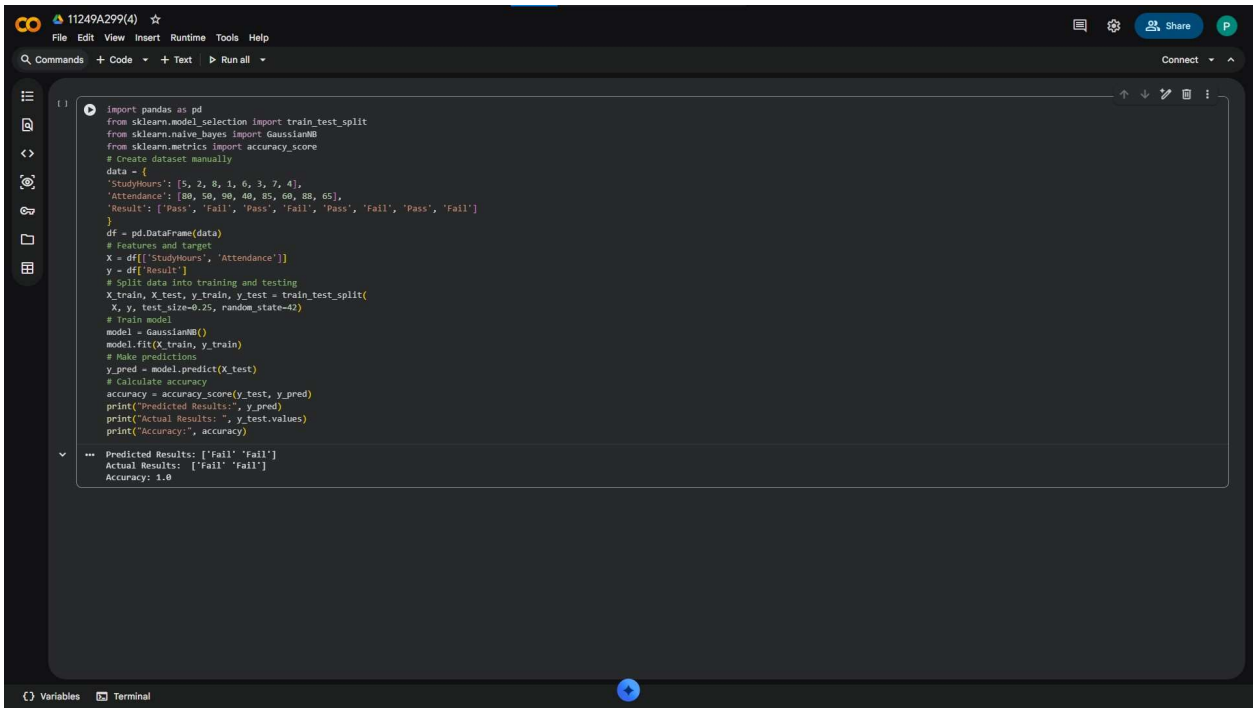

```python
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score
# Create dataset manually
data = {
'StudyHours': [5, 2, 8, 1, 6, 3, 7, 4],
'Attendance': [80, 50, 90, 40, 85, 60, 88, 65],
'Result': ['Pass', 'Fail', 'Pass', 'Fail', 'Pass', 'Fail', 'Pass', 'Fail']
}
df = pd.DataFrame(data)
# Features and target
X = df[['StudyHours', 'Attendance']]
y = df['Result']
# Split data into training and testing
X_train, X_test, y_train, y_test = train_test_split(
 X, y, test_size=0.25, random_state=42)
# Train model
model = GaussianNB()
model.fit(X_train, y_train)
# Make predictions
y_pred = model.predict(X_test)
# Calculate accuracy
accuracy = accuracy_score(y_test, y_pred)
print("Predicted Results:", y_pred)
print("Actual Results: ", y_test.values)
print("Accuracy:", accuracy)
```

```
Predicted Results: ['Fail' 'Fail']
Actual Results:  ['Fail' 'Fail']
Accuracy: 1.0
```

Result: The Classifier was implemented successfully.

Exp

5:-Implement Regularised logistic regression

Aim: To implement regularized logistic regression.

```python
# 1. Import Libraries
import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import load_breast_cancer
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score

# 2. Load Dataset
data = load_breast_cancer()
X = data.data
y = data.target

# 3. Split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

# 4. Scaling
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# 5. Normal Logistic Regression (Very Large C)
normal_model = LogisticRegression(C=1e6, penalty='l2', solver='liblinear')
normal_model.fit(X_train_scaled, y_train)

normal_pred = normal_model.predict(X_test_scaled)
normal_acc = accuracy_score(y_test, normal_pred)

# 6. Regularized Logistic Regression
regularized_model = LogisticRegression(C=0.1, penalty='l2', solver='liblinear')
regularized_model.fit(X_train_scaled, y_train)

reg_pred = regularized_model.predict(X_test_scaled)
reg_acc = accuracy_score(y_test, reg_pred)

# 7. Print Results
print("Normal Logistic Accuracy:", normal_acc * 100)
print("Regularized Logistic Accuracy:", reg_acc * 100)

# 8. Compare Coefficients
plt.figure()
plt.plot(normal_model.coef_.flatten(), label='Normal Logistic')
plt.plot(regularized_model.coef_.flatten(), label='Regularized Logistic')
plt.legend()
```

```python
# 8. Compare Coefficients
plt.figure()
plt.plot(normal_model.coef_.flatten(), label='Normal Logistic')
plt.plot(regularized_model.coef_.flatten(), label='Regularized Logistic')
plt.legend()
plt.title("Coefficient Comparison")
plt.xlabel("Feature Index")
plt.ylabel("Coefficient Value")
plt.show()

# 9. Accuracy Comparison Bar Graph
plt.figure()
models = ['Normal Logistic', 'Regularized Logistic']
accuracies = [normal_acc, reg_acc]

plt.bar(models, accuracies)
plt.title("Accuracy Comparison")
plt.ylabel("Accuracy")
plt.show()
```

```
Normal Logistic Accuracy: 93.85964912280701
Regularized Logistic Accuracy: 99.12280701754386
```

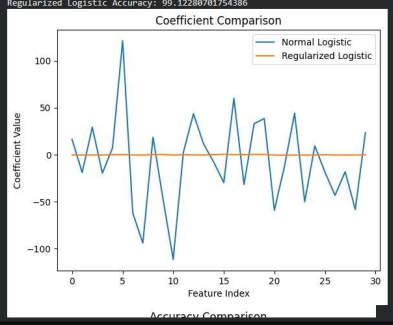

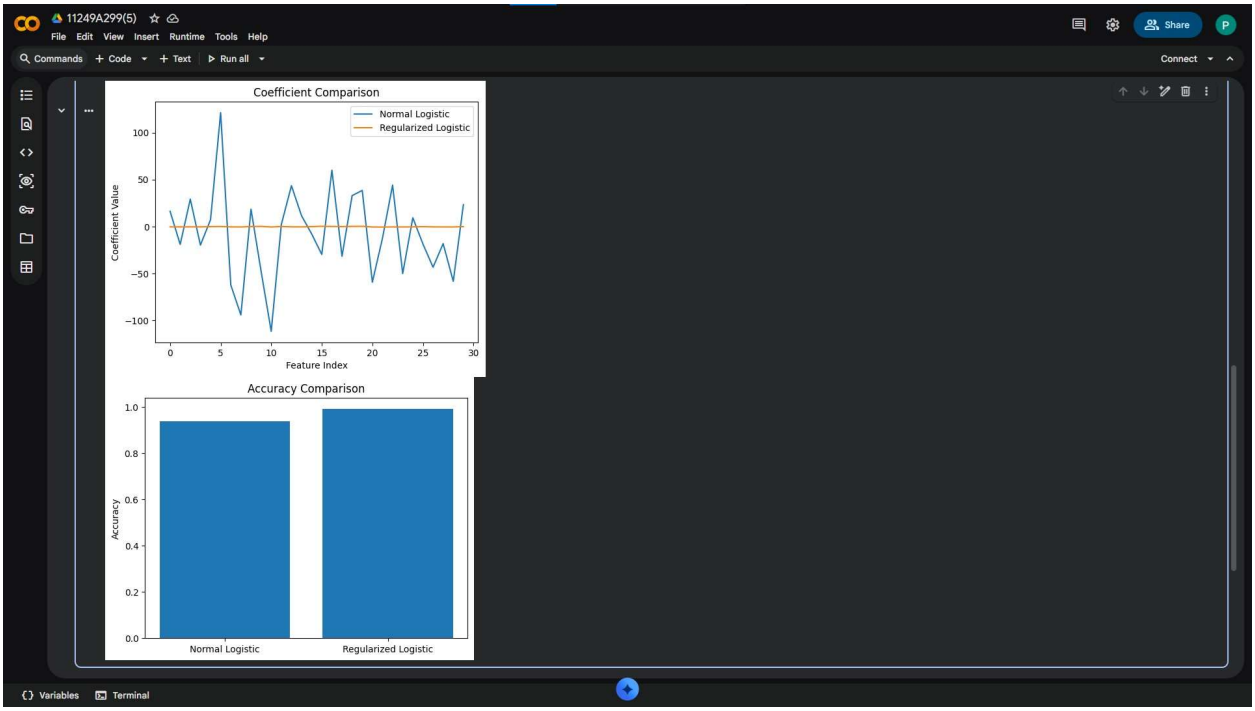

Result: The logistic regression model was implemented successfully.

6 : Build Models using different Ensembling Techniques

Aim: To implement ensembling techniques.

Exp

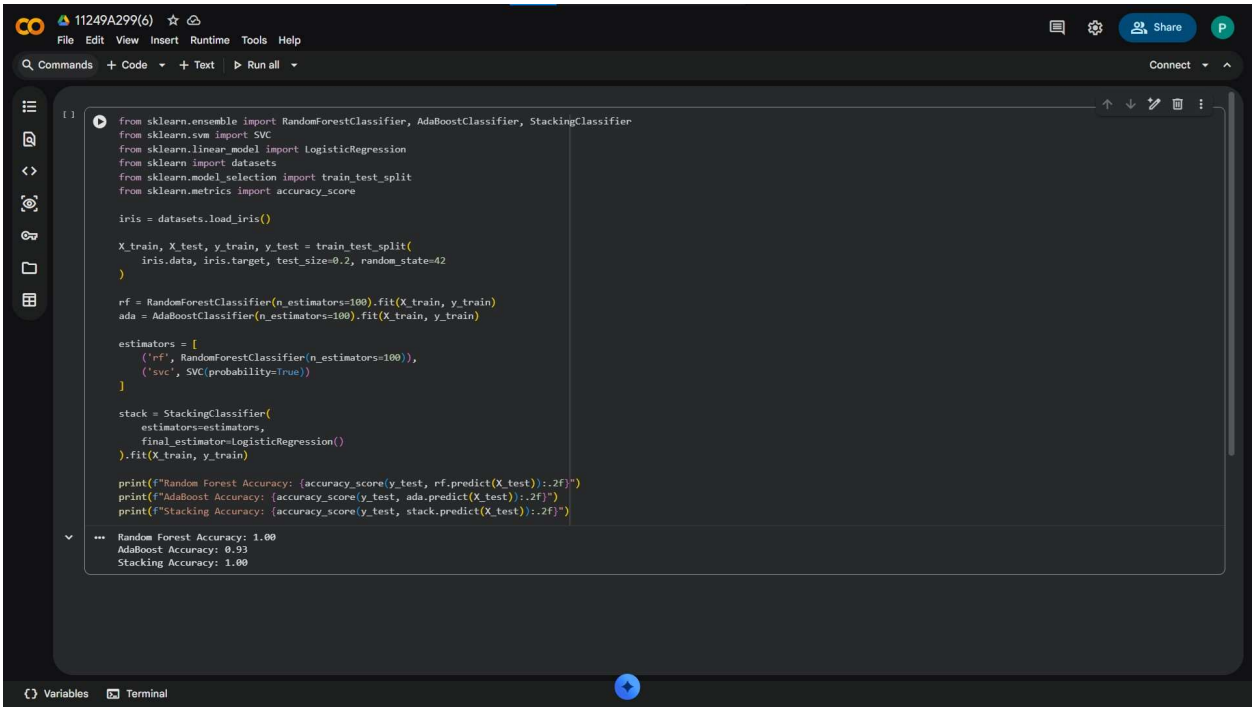

```python
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier, StackingClassifier
from sklearn.svm import SVC
from sklearn.linear_model import LogisticRegression
from sklearn import datasets
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score

iris = datasets.load_iris()

X_train, X_test, y_train, y_test = train_test_split(
    iris.data, iris.target, test_size=0.2, random_state=42
)

rf = RandomForestClassifier(n_estimators=100).fit(X_train, y_train)
ada = AdaBoostClassifier(n_estimators=100).fit(X_train, y_train)

estimators = [
    ('rf', RandomForestClassifier(n_estimators=100)),
    ('svc', SVC(probability=True))
]

stack = StackingClassifier(
    estimators=estimators,
    final_estimator=LogisticRegression()
).fit(X_train, y_train)

print(f"Random Forest Accuracy: {accuracy_score(y_test, rf.predict(X_test)):.2f}")
print(f"AdaBoost Accuracy: {accuracy_score(y_test, ada.predict(X_test)):.2f}")
print(f"Stacking Accuracy: {accuracy_score(y_test, stack.predict(X_test)):.2f}")
```

```
Random Forest Accuracy: 1.00
AdaBoost Accuracy: 0.93
Stacking Accuracy: 1.00
```

Result: The ensembling techniques were implemented successfully.

Exp 7 : Build Models using Decision Trees

Aim: To implement decision tree models.

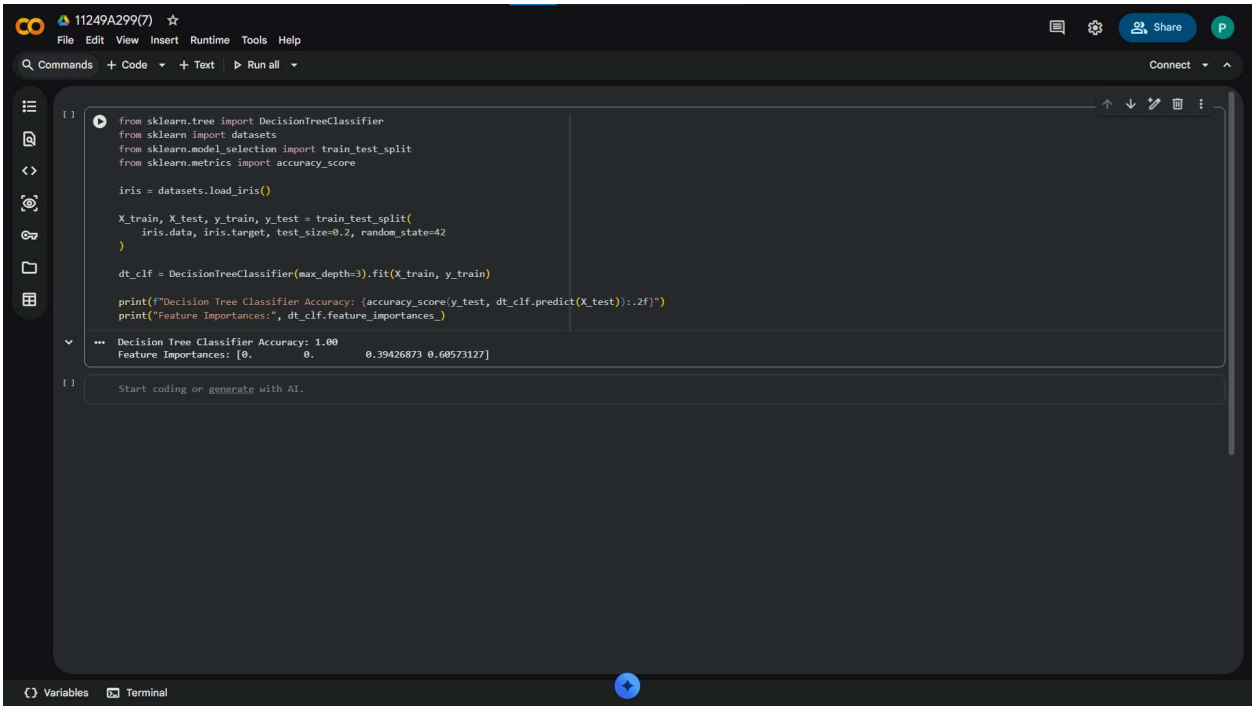

```
from sklearn.tree import DecisionTreeClassifier
from sklearn import datasets
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score

iris = datasets.load_iris()

X_train, X_test, y_train, y_test = train_test_split(
    iris.data, iris.target, test_size=0.2, random_state=42
)

dt_clf = DecisionTreeClassifier(max_depth=3).fit(X_train, y_train)

print(f"Decision Tree Classifier Accuracy: {accuracy_score(y_test, dt_clf.predict(X_test)):.2f}")
print("Feature Importances:", dt_clf.feature_importances_)
```

```
Decision Tree Classifier Accuracy: 1.00
Feature Importances: [0.         0.         0.39426873 0.60573127]
```

Result: The decision tree models were implemented successfully.

Exp 8: Support Vector Machines (SVM) with various Kernels

Aim: To implement a Support Vector Machine.

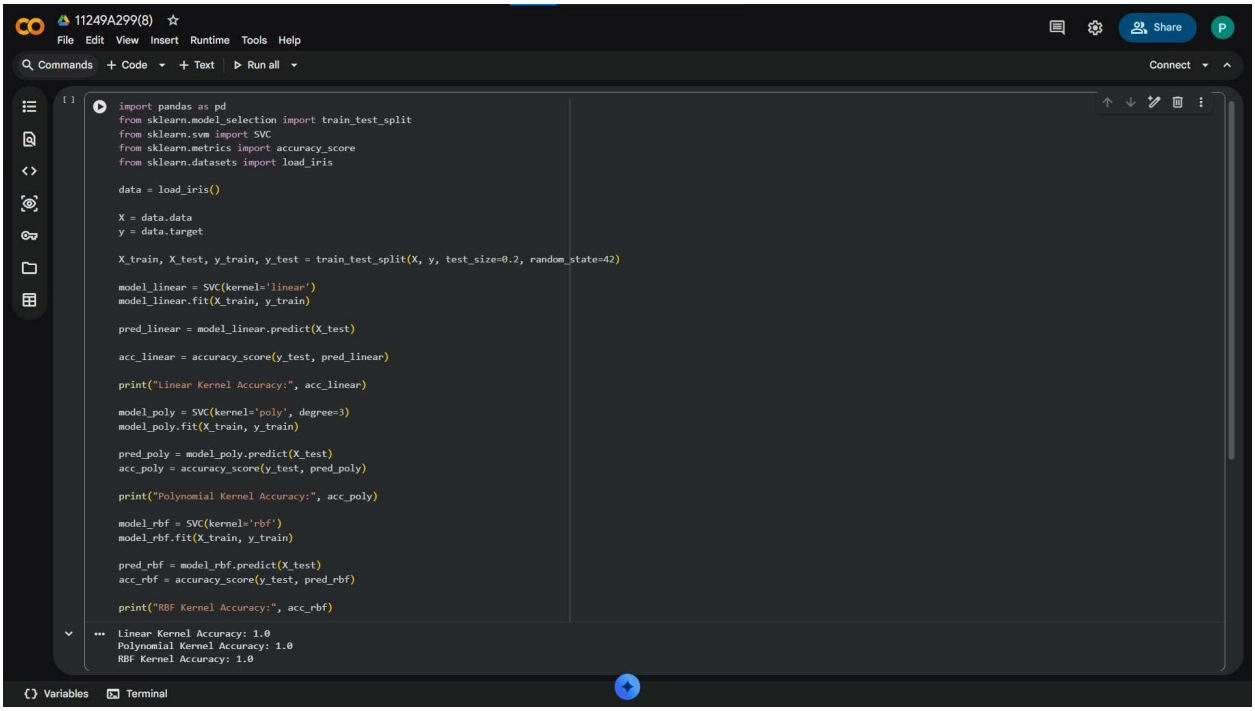

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score
from sklearn.datasets import load_iris

data = load_iris()

X = data.data
y = data.target

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

model_linear = SVC(kernel='linear')
model_linear.fit(X_train, y_train)

pred_linear = model_linear.predict(X_test)

acc_linear = accuracy_score(y_test, pred_linear)

print("Linear Kernel Accuracy:", acc_linear)

model_poly = SVC(kernel='poly', degree=3)
model_poly.fit(X_train, y_train)

pred_poly = model_poly.predict(X_test)
acc_poly = accuracy_score(y_test, pred_poly)

print("Polynomial Kernel Accuracy:", acc_poly)

model_rbf = SVC(kernel='rbf')
model_rbf.fit(X_train, y_train)

pred_rbf = model_rbf.predict(X_test)
acc_rbf = accuracy_score(y_test, pred_rbf)

print("RBF Kernel Accuracy:", acc_rbf)
```

```
Linear Kernel Accuracy: 1.0
Polynomial Kernel Accuracy: 1.0
RBF Kernel Accuracy: 1.0
```

Result: The Support Vector Machines with Kernels was implemented successfully

Exp 9:-Implement KNN Classification

Aim: To implement K-Nearest Neighbors algorithm

```python
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

data = pd.read_csv("/content/mobile price.csv")

X = data.drop("price_range", axis=1)
y = data["price_range"]

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

model = KNeighborsClassifier()
model.fit(X_train, y_train)

y_pred = model.predict(X_test)

print("Accuracy:", accuracy_score(y_test, y_pred))
print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred))
```

```
Accuracy: 0.53
[[76 24  5  0]
 [25 41 22  3]
 [ 6 41 34 11]
 [ 2 13 36 61]]
              precision    recall  f1-score   support

           0       0.70      0.72      0.71       105
           1       0.34      0.45      0.39        91
           2       0.35      0.37      0.36        92
           3       0.81      0.54      0.65       112

    accuracy                           0.53       400
   macro avg       0.55      0.52      0.53       400
weighted avg       0.57      0.53      0.54       400
```

Result: The KNN Classification model was implemented successfully.

Exp 10 : K- Means Clustering with PCA and Elbow Method

Aim: To implement k-means clustering.

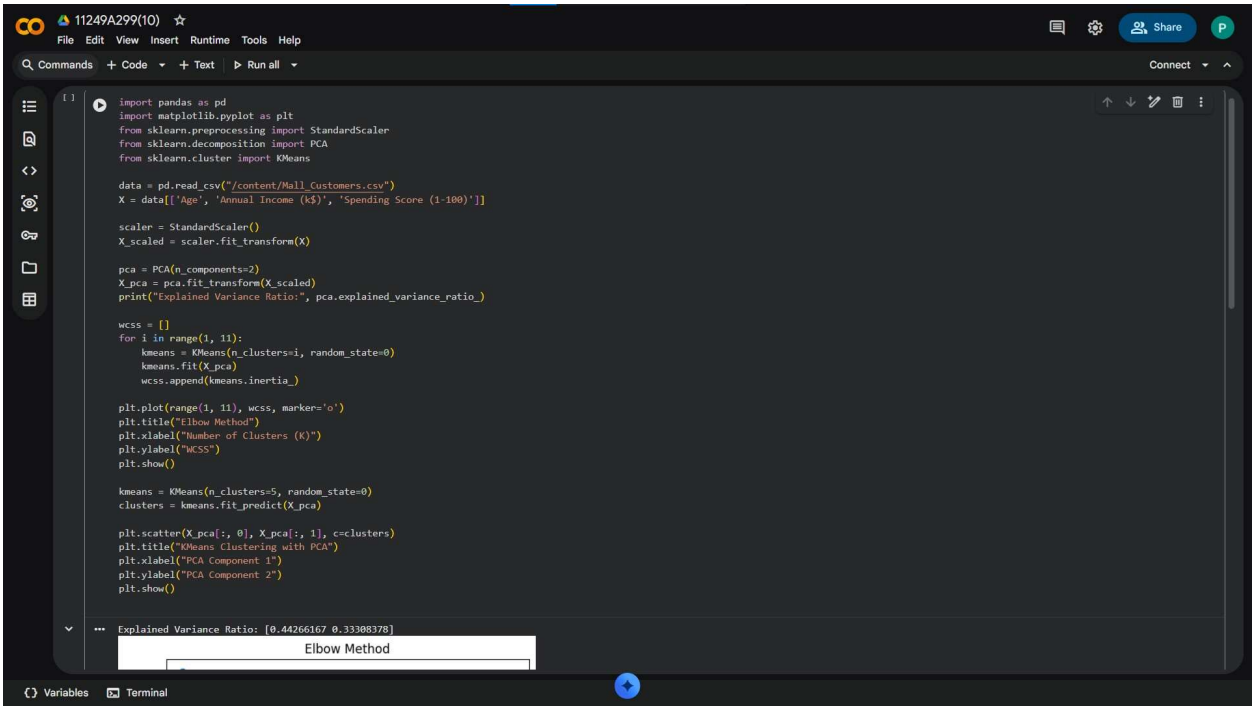

```python
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.cluster import KMeans

data = pd.read_csv("/content/Mall_Customers.csv")
X = data[['Age', 'Annual Income (k$)', 'Spending Score (1-100)']]

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

pca = PCA(n_components=2)
X_pca = pca.fit_transform(X_scaled)
print("Explained Variance Ratio:", pca.explained_variance_ratio_)

wcss = []
for i in range(1, 11):
    kmeans = KMeans(n_clusters=i, random_state=0)
    kmeans.fit(X_pca)
    wcss.append(kmeans.inertia_)

plt.plot(range(1, 11), wcss, marker='o')
plt.title("Elbow Method")
plt.xlabel("Number of Clusters (K)")
plt.ylabel("WCSS")
plt.show()

kmeans = KMeans(n_clusters=5, random_state=0)
clusters = kmeans.fit_predict(X_pca)

plt.scatter(X_pca[:, 0], X_pca[:, 1], c=clusters)
plt.title("KMeans Clustering with PCA")
plt.xlabel("PCA Component 1")
plt.ylabel("PCA Component 2")
plt.show()
```

```
Explained Variance Ratio: [0.44266167 0.33308378]
```

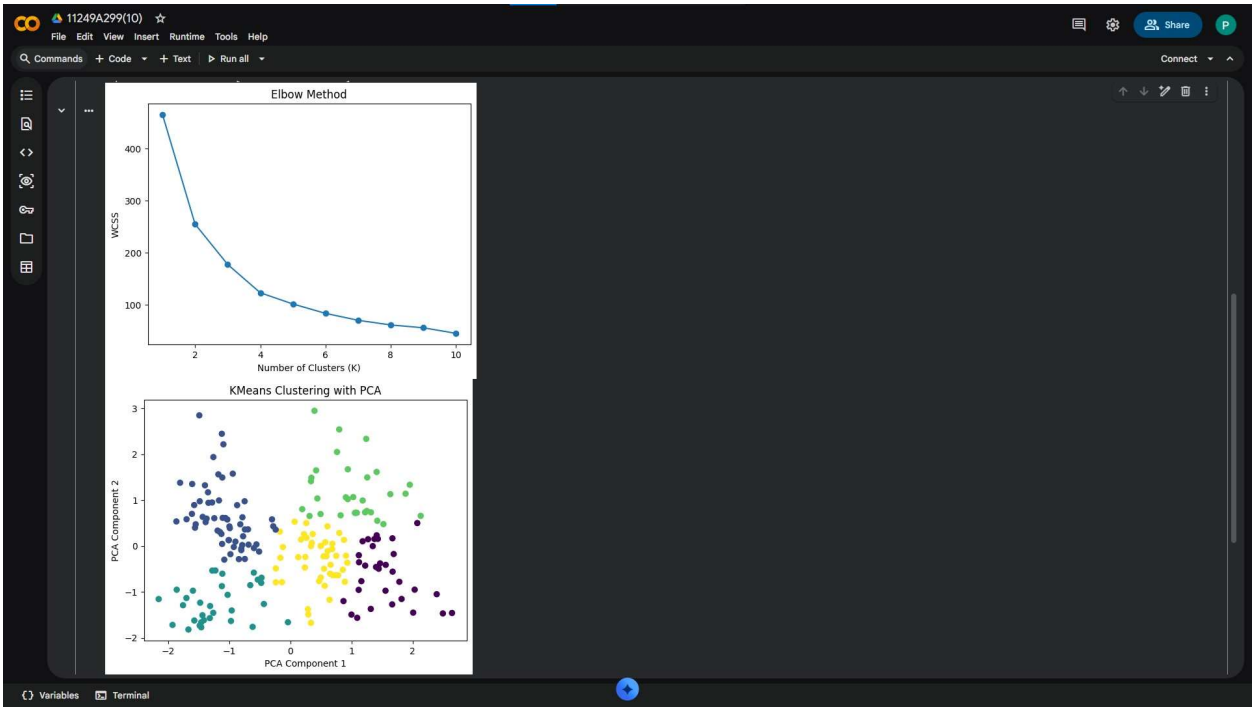

Result: K- Means Clustering with PCA and Elbow Method was performed successfully.